

Benchmarking Symbolic Execution Using Constraint Problems - Initial Results

Sahil Verma, IIT Kanpur

Roland Yap, National University of Singapore

November 4, 2019

Overview

Motivation

Constraint Satisfaction Problems (CSP)

Transforming CSP to C

Experiments

Conclusion

Motivation

- Symbolic execution engines have played phenomenal role in program analysis
- We aim to improve symbolic execution via standardized benchmarking
- We are motivated by how “hard benchmarks” from SAT competition has driven the development of SAT solvers.

High-level idea

- We propose synthetic benchmarks which exercise the core reasoning techniques
- We propose to use discrete combinatorial problems in form of constraint satisfaction problems (CSP)
- We transform the CSPs' into C code using various possible transformations to construct the benchmarks

CSP background

- A Constraint Satisfaction Problem(CSP) P is a pair (X, C) where X is a set of n variables x_1, \dots, x_n and C a set of e constraints c_1, \dots, c_e
- In a finite domain CSP, variables $x \in X$ take values from its domain $D(x)$ which is a finite set of values
- The general form of a finite domain CSP is NP-complete

Constraint forms

- Constraints can either be defined implicitly or in form of explicit values
- When constraints are implicitly defined in form of linear arithmetic relations over variables in X , they are termed as intensional constraints

$\langle \text{constraint} \rangle \text{equal}([\%0], \text{sum}([\%1], [\%2])) \langle \text{constraint} \rangle$
 $\langle \text{args} \rangle x[0] \ x[0] \ x[1] \ \langle / \text{args} \rangle$
 $\langle \text{args} \rangle x[1] \ x[1] \ x[2] \ \langle / \text{args} \rangle$

Intensional Constraint

Constraint forms

- Constraints can be defined explicitly by providing a list of values some set of variables satisfy (positive tables) or do not satisfy (negative tables)

$\langle \textit{positive} \rangle (0, 0, 0) (0, 1, 0) \langle / \textit{positive} \rangle$
 $\langle \textit{args} \rangle x[0] \ x[1] \ x[2] \ \langle / \textit{args} \rangle$
 $\langle \textit{args} \rangle x[3] \ x[4] \ x[5] \ \langle / \textit{args} \rangle$

Extensional Constraint

Transforming CSP to C

- The finite domain variables of the CSP correspond to integer variables in the program P'
- The C variables are made symbolic
- The constraint relations are encoded into conditional statements in the program (if, else) or as 'assume' statement

Transforming CSP to C

- We generate several equisatisfiable transformations for a CSP
- The transformations differ in the C constructs, the constraint grouping aggressiveness and the operators used for grouping

Type	Versions	Construct	Operator	Grouped
Extensional	1	if	<i>logical</i>	<i>no</i>
	2	if	<i>logical</i>	<i>yes</i>
	3	if	<i>bitwise</i>	<i>no</i>
	4	<i>assume</i>	<i>bitwise</i>	<i>yes</i>
	...			
Intensional	1	if	<i>NOP</i>	<i>no</i>
	2	<i>assume</i>	<i>logical</i>	<i>yes</i>
	3	<i>assume</i>	<i>bitwise</i>	<i>no</i>
	4	<i>assume</i>	<i>bitwise</i>	<i>yes</i>
	...			

Distinguishing SAT and UNSAT CSPs

- When all constraints have been transformed, we place an 'assert(0)' at the end of the program to trigger a failure flag
- If the CSP P is satisfiable, the program terminates with an assertion failure, and if it is unsatisfiable, the statement is unreachable.
- We did not find any contradictions in our experiments

Experimental Setup

- We generate 12 extensional and 10 intensional versions
- We experiment with 6 classes, 3 from intensional and 3 from extensional category totalling 119 programs
- We use KLEE (with both STP and Z3 as solvers), Tracer-X, LLBMC as execution engines. Our baseline was AbsCon, a common tool for solving CSPs

Experimental Setup

- We use the symbolic execution engine's runtime for a transformed program P' as a measure of the tool's robustness over transformations
- We use the number of timeouts for benchmarks with a size parameter as a measure of tool's scalability with increasing size of programs

Results




Figure: Robustness of different tools for an extensional class benchmark (AIM-100 satisfiable and unsatisfiable)

Results




Figure: Scalability of different tools for an intensional class benchmark (CostasArrays)

Conclusions

- The results show that there certainly is a significant amount of fragility in the tools which could be addressed
- The results also show that different solvers behave differently when used in KLEE, but in an unpredictable fashion
- Tracer-X which uses interpolation to remove execution paths was faster for some benchmarks and slower for some, which points towards prediction techniques which would help direct its usage